

# Kehan Guo

[kguo2@nd.edu](mailto:kguo2@nd.edu)   [kehanguo2.github.io](https://kehanguo2.github.io)   [github.com/KehanGuo2](https://github.com/KehanGuo2)   [Google Scholar](#)

## SUMMARY

Ph.D. candidate at Notre Dame working on RL post-training and reward design for LLM agents. One question drives the work: when a model is optimized against a proxy — a reward model, a verifier, a frozen environment — where does the pressure actually go? Two answers are under review — a zero-cost reward interface in flow matching’s noise–data coupling [5] and a live evaluation harness for tool-using agents [6] — both grown from community-standard benchmarks I created and co-led [1, 3]. Joining Apple as a research intern, June 2026; previously Amazon AWS AI and ByteDance.

## EDUCATION

### University of Notre Dame

Aug 2022 – May 2027 (expected)

*Ph.D. in Computer Science · Advisor: Prof. Xiangliang Zhang (MINE Lab)*

### Boston University

2020 – 2022

*M.S. in Computer Science*

## RESEARCH EXPERIENCE

### 🍏 Apple

Jun 2026 – present

*Research Intern · Cupertino, CA*

- Research area: reinforcement-learning post-training for large language models.

### ByteDance

Apr 2026 – Jun 2026

*Research Scientist Intern · San Jose, CA*

- Built BI-Claw, an internal tool-using LLM agent for business-integrity risk investigation, orchestrating 20+ domain skills over internal data services.
- Designed a self-evolving skill pipeline — skills seeded from static evaluation, evolved from real user queries with LLM-as-judge feedback — raising task success from 62% to 70% on the team’s 145-case internal benchmark.
- Built the agentic evaluation harness behind that number — the industrial counterpart of [6].

### 🍎 Amazon AWS AI

May 2025 – Aug 2025

*Applied Scientist Intern · New York, NY*

- Designed **Social-Choice GRPO**, a reward shaping framework that converts sparse binary RL signals into dense rank-based feedback via social choice aggregation across policy rollouts, reducing gradient variance on long-horizon reasoning chains.
- Built a model-agnostic critique layer fusing rule-based verification with calibrated model confidence, enabling stable policy gradient updates on ambiguous, multi-step reasoning paths.
- Improved accuracy and sample efficiency over PPO baselines across math, coding, and logical reasoning tasks.

### 🏰 University of Notre Dame — MINE Lab

Aug 2022 – Present

*Research Assistant*

*Reward-Aligned Generation — Reward Transport [5] (under review)*

- *Can alignment signal enter generative training through interfaces we already have?* Showed the noise–data coupling in flow matching — long treated as a computational convenience — is a zero-cost interface for property alignment, and characterized exactly when the interface exists.

*Live Agentic Evaluation — ChemElucid-Live [6] (under review)*

- *What do benchmark scores miss once the world moves?* Built a live evaluation harness measuring how tool-using agents acquire, verify, and act on information — the live complement to static benchmarks.

*Multimodal Reasoning — MolPuzzle (NeurIPS 2024 [Spotlight]) [1]*

- *Can frontier models do expert perception, not just expert recall?* Created MolPuzzle, a multimodal benchmark for molecular structure elucidation from spectra, targeting the perception–reasoning gap.

*LLM Evaluation for Science — ChemLLMBench (NeurIPS 2023) [3]*

- *What can LLMs actually do in chemistry?* Co-led an early comprehensive benchmark of LLM chemistry capabilities across eight tasks.

*Selected Collaborations*

- Co-authored TRACE (runtime preference enforcement for coding agents, 2026); contributed ideation and pipeline design to AdaReasoner (adaptive inference-time reasoning, NeurIPS 2025 Spotlight) and ChemOrch (74-tool agent orchestration, NeurIPS 2025).

## SELECTED PUBLICATIONS

---

\* equal contribution. Full list at [Google Scholar](#).

- [1] Can LLMs Solve Molecule Puzzles? A Multimodal Benchmark for Molecular Structure Elucidation.  
**K. Guo**, B. Nan, Y. Zhou, et al. *NeurIPS 2024 [Spotlight]*
- [2] Evaluating Large Language Models in Scientific Discovery.  
Z. Song\*, J. Lu\*, Y. Du, B. Yu, T. M. Pruyn, Y. Huang, **K. Guo\***, et al. *arXiv 2025*
- [3] What Can Large Language Models Do in Chemistry? A Comprehensive Benchmark on Eight Tasks.  
T. Guo\*, **K. Guo\***, B. Nan, et al. *NeurIPS 2023*
- [4] Artificial Intelligence in Spectroscopy: Advancing Chemistry from Prediction to Generation and Beyond.  
**K. Guo**, Y. Shen, et al. *IJCAI 2025 [Survey Track]*

## Preprints / Under Review

- [5] Reward Transport: Property Control in Flow Matching via Noise-Space Alignment.  
**K. Guo**, Y. Shen, Y. Zhou, Y. Huang, C. Gao, S. Du, X. Zhang. *Under review [Code]*
- [6] ChemElucid-Live: Live Evaluation of Tool-Using Scientific Agents.  
**K. Guo**, et al. *Under review*

## TECHNICAL SKILLS

---

|                               |                                                                                               |
|-------------------------------|-----------------------------------------------------------------------------------------------|
| <b>Languages</b>              | Python, C++, SQL, Bash                                                                        |
| <b>Training</b>               | PyTorch, JAX, Megatron-LM, DeepSpeed ZeRO, Ray, VeRL, GRPO/PPO/DPO pipelines, reward modeling |
| <b>Inference &amp; Agents</b> | vLLM, HF Transformers, function-calling, Neo4j                                                |
| <b>Dev Tools</b>              | AI-assisted development: Claude Code, Codex                                                   |
| <b>Scientific</b>             | RDKit, PubChem, spectroscopy processing                                                       |

## AWARDS & RECOGNITION

---

NeurIPS 2024 Spotlight — *MolPuzzle* (first author)  
NeurIPS 2025 Spotlight — *AdaReasoner* (co-author)  
OpenAI Researcher Access Program (2024)

## PROFESSIONAL SERVICE

---

**Reviewer:** NeurIPS (2024–25), ICLR (2025–26), ICML, AAAI, IJCAI, KDD, WWW, ACL Rolling Review (ACL/EMNLP)